

Introduction to Machine Learning

Varun Chandola <chandola@buffalo.edu>

June 30, 2014

Outline

- 1 Supervised Learning
- 2 Grouping Learning Algorithms
- 3 Maximum Margin Classifiers
- 4 Support Vector Machines
- 5 More About Kernels

Outline

- 1 **Supervised Learning**
 - Learning Parameters
 - Risk Minimization for Parameter Learning
- 2 **Grouping Learning Algorithms**
 - 0-1 Loss
 - Approximation to 0-1 Loss
- 3 **Maximum Margin Classifiers**
 - Boundary with Maximum Margin
- 4 **Support Vector Machines**
 - SVM Learning
 - SVM for Non-Separable Case
 - Optimization Constraints
- 5 **More About Kernels**
 - Gaussian Kernel

Supervised Learning

- **Classification Task:** Use a model (or hypothesis) to assign a label (or value) to a previously unseen instance, \mathbf{x}
- **Learning Task:** *Learn* the hypothesis or model parameters using training data $\langle \mathbf{x}_i, y_i \rangle_{i=1}^N$

Ideal Learning Algorithm

- 1 Fast
- 2 Generalizable

Learning Parameters Via Loss Minimization

- **Notion of Loss** - how badly does the model (\mathcal{M}) perform on a \mathbf{x} if y is the correct output?
- What can the loss function be ($loss(\mathbf{x}, y; \mathcal{M})$)?
- How do you measure this loss?
- If we had all possible realizations of \mathbf{x} and corresponding y
- If we have a probability distribution defined over \mathbf{x}, y

$$\mathbb{E}_{p(\mathbf{x}, y)}[loss(\mathbf{x}, y; \mathcal{M})]$$

- **Risk minimization**
- Usually we do not have any of these

Learning Parameters Via Loss Minimization

- **Notion of Loss** - how badly does the model (\mathcal{M}) perform on a \mathbf{x} if y is the correct output?
- What can the loss function be ($loss(\mathbf{x}, y; \mathcal{M})$)?
- How do you measure this loss?
- If we had all possible realizations of \mathbf{x} and corresponding y
- If we have a probability distribution defined over \mathbf{x}, y

$$\mathbb{E}_{p(\mathbf{x}, y)}[loss(\mathbf{x}, y; \mathcal{M})]$$

- **Risk minimization**
- Usually we do not have any of these

Empirical Risk Minimization

- What we do have is an *empirical distribution* (from the training data)!
- **Empirical Risk**

$$\mathbb{E}_{\tilde{p}(\mathbf{x}, y)}[\text{loss}(\mathbf{x}, y; \mathcal{M})] = \frac{1}{N} \sum_{i=1}^N [\text{loss}(\mathbf{x}_i, y_i; \mathcal{M})]$$

- Find model parameters (or hypothesis) that minimizes the empirical risk

$$\arg \min_{\mathcal{M}} \mathbb{E}_{\tilde{p}(\mathbf{x}, y)}[\text{loss}(\mathbf{x}, y; \mathcal{M})] = \arg \min_{\mathcal{M}} \frac{1}{N} \sum_{i=1}^N [\text{loss}(\mathbf{x}_i, y_i; \mathcal{M})]$$

Structural Risk Minimization

- What if ERM gives multiple solution with “equal” risk?
 - Which model to choose?
 - Choose the one with least complexity (Akaike Information Criterion, Minimum Description Length principle, **Structural Risk Minimization** (SRM))
 - Ensures better generalizability
- SRM provides a *trade-off* between complexity of an algorithm (**VC Dimension**) and quality of *fit* on the training data (**empirical error**)
 - Low VC dimension \Rightarrow Low complexity \Rightarrow Better generalizability

Structural Risk Minimization

- What if ERM gives multiple solution with “equal” risk?
 - Which model to choose?
 - Choose the one with least complexity (Akaike Information Criterion, Minimum Description Length principle, **Structural Risk Minimization** (SRM))
 - Ensures better generalizability
- SRM provides a *trade-off* between complexity of an algorithm (**VC Dimension**) and quality of *fit* on the training data (**empirical error**)
 - Low VC dimension \Rightarrow Low complexity \Rightarrow Better generalizability

Structural Risk Minimization

- What if ERM gives multiple solution with “equal” risk?
 - Which model to choose?
 - Choose the one with least complexity (Akaike Information Criterion, Minimum Description Length principle, **Structural Risk Minimization** (SRM))
 - Ensures better generalizability
- SRM provides a *trade-off* between complexity of an algorithm (**VC Dimension**) and quality of *fit* on the training data (**empirical error**)
 - Low VC dimension \Rightarrow Low complexity \Rightarrow Better generalizability

Regularized Risk Minimization

- *Regularize* the complexity of the model

$$\arg \min_{\mathcal{M}} \frac{1}{N} \sum_{i=1}^N [\text{loss}(\mathbf{x}_i, y_i; \mathcal{M})] + R(\mathcal{M})$$

Outline

- 1 Supervised Learning
 - Learning Parameters
 - Risk Minimization for Parameter Learning
- 2 Grouping Learning Algorithms
 - 0-1 Loss
 - Approximation to 0-1 Loss
- 3 Maximum Margin Classifiers
 - Boundary with Maximum Margin
- 4 Support Vector Machines
 - SVM Learning
 - SVM for Non-Separable Case
 - Optimization Constraints
- 5 More About Kernels
 - Gaussian Kernel

How do classification algorithms differ from each other?

- What is the output?
 - Probability ($P(y|\mathbf{x})$) - Bayesian Classifier
 - Score ($s(\mathbf{x})$) with a threshold? - Support Vector Machines, Decision Trees
- How is $P(y|\mathbf{x})$ computed?
 - From $P(\mathbf{x}|y)$ (Generative) - Naive Bayes Classifier
 - Directly (Discriminative) - Logistic Regression
- What is the loss function?
 - 0-1 loss (Not easy to regularize ☹)
 - Log loss (Logistic Regression, CRF, Max ent models)
 - Hinge Loss (SVM)
- What is the regularizer?

Jumping to Linear Classifiers

- Calculate $\mathbf{w}^\top \mathbf{x} + b$ - linear combination of features
- Learning can be posed as a general optimization problem

$$\min_{\mathbf{w}, b} L(\mathbf{w}, b) = \min_{\mathbf{w}, b} \sum_{n=1}^N \mathbb{I}(y_n(\mathbf{w}^\top \mathbf{x}_n + b) < 0) + \lambda R(\mathbf{w}, b)$$


- \mathbb{I} is an **indicator function** (1 if $(.)$ is 0, 0 otherwise)
- Objective function = **Loss function** + λ **Regularizer**
- Objective function wants to **fit training data well** and **have simpler solution**

0-1 Loss is Hard to Optimize

- Combinatorial optimization problem
- NP-hard
- No polynomial time algorithm
- Loss function is non-smooth, non-convex
- Small changes in \mathbf{w}, b can change the loss by lot


Approximations to 0-1 Loss

- Different linear classifiers use different approximations to 0-1 loss
 - Also known as *surrogate loss functions*




Approximations to 0-1 Loss

- Different linear classifiers use different approximations to 0-1 loss
 - Also known as *surrogate loss functions*



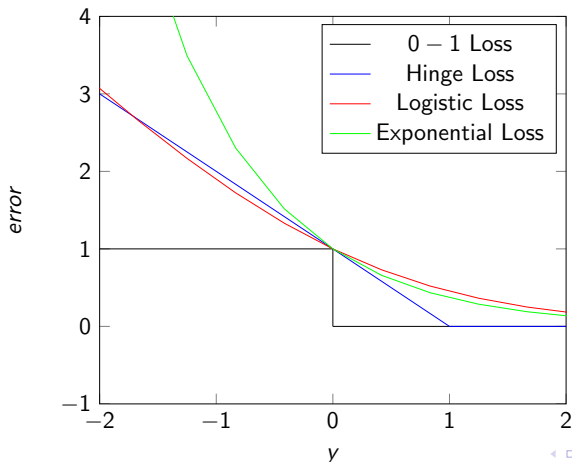
Approximations to 0-1 Loss

- Different linear classifiers use different approximations to 0-1 loss
 - Also known as *surrogate loss functions*




Approximations to 0-1 Loss

- Different linear classifiers use different approximations to 0-1 loss
 - Also known as *surrogate loss functions*



Approximations to 0-1 Loss

- Different linear classifiers use different approximations to 0-1 loss
 - Also known as *surrogate loss functions*



Role of Regularizers

- Recall the optimization problem for linear classification

$$\min_{\mathbf{w}, b} L(\mathbf{w}, b) = \min_{\mathbf{w}, b} \sum_{n=1}^N \mathbb{I}(y_n(\mathbf{w}^\top \mathbf{x}_n + b) < 0) + \lambda R(\mathbf{w}, b)$$

- What is the role of the regularizer term?
 - Ensure simplicity
- Ideally we want most entries of \mathbf{w} to be zero
- Why?
- Desired minimization

$$R(\mathbf{w}, b) = \sum_{d=1}^D \mathbb{I}(w_d \neq 0)$$

- NP Hard

Role of Regularizers

- Recall the optimization problem for linear classification

$$\min_{\mathbf{w}, b} L(\mathbf{w}, b) = \min_{\mathbf{w}, b} \sum_{n=1}^N \mathbb{I}(y_n(\mathbf{w}^\top \mathbf{x}_n + b) < 0) + \lambda R(\mathbf{w}, b)$$

- What is the role of the regularizer term?
 - Ensure simplicity
- Ideally we want most entries of \mathbf{w} to be zero
- Why?
- Desired minimization

$$R(\mathbf{w}, b) = \sum_{d=1}^D \mathbb{I}(w_d \neq 0)$$

- NP Hard

Approximate Regularization

- **Norm based regularization**

- l_2 squared norm

$$\|\mathbf{w}\|_2^2 = \sum_{d=1}^D w_d^2$$

- l_1 norm

$$\|\mathbf{w}\|_1 = \sum_{d=1}^D |w_d|$$


- l_p norm

$$\|\mathbf{w}\|_p = \left(\sum_{d=1}^D w_d^p \right)^{1/p}$$

- Norm becomes non-convex for $p < 1$
- l_1 norm gives best results
- l_2 norm is easiest to deal with

Linear Hyperplane

- Separates a D -dimensional space into two half-spaces
- Defined by $\mathbf{w} \in \mathbb{R}^D$
 - *Orthogonal* to the hyperplane
 - This \mathbf{w} goes through the origin
 - How do you check if a point lies “above” or “below” \mathbf{w} ?
 - What happens for points **on** \mathbf{w} ?



Make hyperplane not go through origin

- Add a bias b
 - $b > 0$ - move along \mathbf{w}
 - $b < 0$ - move opposite to \mathbf{w}
- How to check if point lies above or below \mathbf{w} ?
 - If $\mathbf{w}^\top \mathbf{x} + b > 0$ then \mathbf{x} is *above*
 - Else, *below*


Outline

- 1 Supervised Learning
 - Learning Parameters
 - Risk Minimization for Parameter Learning
- 2 Grouping Learning Algorithms
 - 0-1 Loss
 - Approximation to 0-1 Loss
- 3 **Maximum Margin Classifiers**
 - Boundary with Maximum Margin
- 4 Support Vector Machines
 - SVM Learning
 - SVM for Non-Separable Case
 - Optimization Constraints
- 5 More About Kernels
 - Gaussian Kernel

Maximum Margin Classifiers

$$y = \mathbf{w}^T \mathbf{x} + b$$


- Remember the Perceptron!
- If data is linearly separable
 - Perceptron training guarantees learning the decision boundary
- There can be other boundaries
 - Depends on initial value for \mathbf{w}
- But what is the best boundary?



Maximum Margin Classifiers

$$y = \mathbf{w}^T \mathbf{x} + b$$


- Remember the Perceptron!
- If data is linearly separable
 - Perceptron training guarantees learning the decision boundary
- There can be other boundaries
 - Depends on initial value for \mathbf{w}
- But what is the best boundary?



Maximum Margin Classifiers

$$y = \mathbf{w}^T \mathbf{x} + b$$


- Remember the Perceptron!
- If data is linearly separable
 - Perceptron training guarantees learning the decision boundary
- There can be other boundaries
 - Depends on initial value for \mathbf{w}
- **But what is the best boundary?**



Maximum Margin Classifiers

$$y = \mathbf{w}^T \mathbf{x} + b$$


- Remember the Perceptron!
- If data is linearly separable
 - Perceptron training guarantees learning the decision boundary
- There can be other boundaries
 - Depends on initial value for \mathbf{w}
- **But what is the best boundary?**



Maximum Margin Classifiers


$$y = \mathbf{w}^T \mathbf{x} + b$$

- Remember the Perceptron!
- If data is linearly separable
 - Perceptron training guarantees learning the decision boundary
- There can be other boundaries
 - Depends on initial value for \mathbf{w}
- **But what is the best boundary?**



The Notion of Maximum Margin

- If multiple solutions classify the training data perfectly
- Find one which will give the smallest *generalization error*
- Equivalent to choosing the decision surface with **Maximum Margin**




Line as a Decision Surface

- Decision boundary represented by the hyperplane \mathbf{w}
- For binary classification, \mathbf{w} points **towards** the positive class

Decision Rule

$$y = \text{sign}(\mathbf{w}^\top \mathbf{x} + b)$$

- $\mathbf{w}^\top \mathbf{x} + b > 0 \Rightarrow y = +1$
- $\mathbf{w}^\top \mathbf{x} + b < 0 \Rightarrow y = -1$



What is a Margin?

- **Margin** is the distance between an example and the decision line
- Denoted by γ
- For a positive point:

$$\gamma = \frac{\mathbf{w}^\top \mathbf{x} + b}{\|\mathbf{w}\|}$$

- For a negative point:

$$\gamma = -\frac{\mathbf{w}^\top \mathbf{x} + b}{\|\mathbf{w}\|}$$

Functional Interpretation

- Margin **positive** if prediction is **correct**; **negative** if prediction is **incorrect**

Outline

- 1 Supervised Learning
 - Learning Parameters
 - Risk Minimization for Parameter Learning
- 2 Grouping Learning Algorithms
 - 0-1 Loss
 - Approximation to 0-1 Loss
- 3 Maximum Margin Classifiers
 - Boundary with Maximum Margin
- 4 **Support Vector Machines**
 - SVM Learning
 - SVM for Non-Separable Case
 - Optimization Constraints
- 5 More About Kernels
 - Gaussian Kernel

Support Vector Machines

- A hyperplane based classifier defined by \mathbf{w} and b
- Find hyperplane with *maximum separation margin* on the training data
- Assume that data is linearly separable (will relax this later)
 - Zero training error (loss)

SVM Prediction Rule

$$y = \text{sign}(\mathbf{w}^\top \mathbf{x} + b)$$

SVM Learning

- **Input:** Training data $\{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)\}$
- **Objective:** Learn \mathbf{w} and b that maximizes the margin

SVM Learning

- SVM learning task as an optimization problem
- Find \mathbf{w} and b that gives zero training error
- Maximizes the margin $(= \frac{2}{\|\mathbf{w}\|})$
- Same as minimizing $\|\mathbf{w}\|$

Optimization Formulation

$$\begin{aligned} & \underset{\mathbf{w}, b}{\text{minimize}} && \frac{\|\mathbf{w}\|^2}{2} \\ & \text{subject to} && y_n(\mathbf{w}^\top \mathbf{x}_n + b) \geq 1, \quad n = 1, \dots, N. \end{aligned}$$

- **Optimization** with N linear inequality constraint

A Different Interpretation of Margin

- What impact does the margin have on \mathbf{w} ?
- Large margin \Rightarrow Small $\|\mathbf{w}\|$
- Small $\|\mathbf{w}\| \Rightarrow$ regularized/simple solutions
- Simple solutions \Rightarrow Better generalizability (*Occam's Razor*)

Solving the Quadratic Optimization Problem

Optimization Formulation

$$\begin{aligned} & \underset{\mathbf{w}, b}{\text{minimize}} && \frac{\|\mathbf{w}\|^2}{2} \\ & \text{subject to} && y_n(\mathbf{w}^\top \mathbf{x}_n + b) \geq 1, \quad n = 1, \dots, N. \end{aligned}$$

- There is one quantity to minimize and N constraints
- **Primal formulation** - Lagrange Multipliers
- Bring constraints into the objective function

Primal Lagrangian Formulation

$$\begin{aligned} & \underset{\mathbf{w}, b, \alpha}{\text{minimize}} && L_P(\mathbf{w}, b, \alpha) = \frac{\|\mathbf{w}\|^2}{2} + \sum_{n=1}^N \alpha_n (1 - y_n(\mathbf{w}^\top \mathbf{x}_n + b)) \\ & \text{subject to} && \alpha_n \geq 0; \quad n = 1, \dots, N. \end{aligned}$$

Solving the Quadratic Optimization Problem

Optimization Formulation

$$\begin{aligned} & \underset{\mathbf{w}, b}{\text{minimize}} && \frac{\|\mathbf{w}\|^2}{2} \\ & \text{subject to} && y_n(\mathbf{w}^\top \mathbf{x}_n + b) \geq 1, \quad n = 1, \dots, N. \end{aligned}$$

- There is one quantity to minimize and N constraints
- **Primal formulation** - Lagrange Multipliers
- Bring constraints into the objective function

Primal Lagrangian Formulation

$$\begin{aligned} & \underset{\mathbf{w}, b, \alpha}{\text{minimize}} && L_P(\mathbf{w}, b, \alpha) = \frac{\|\mathbf{w}\|^2}{2} + \sum_{n=1}^N \alpha_n (1 - y_n(\mathbf{w}^\top \mathbf{x}_n + b)) \\ & \text{subject to} && \alpha_n \geq 0; \quad n = 1, \dots, N. \end{aligned}$$

More About Quadratic Optimization

- The **Lagrangian** is a lower bound on the original problem
- Find *optimal* values of \mathbf{w} and b , w.r.t. α by setting the derivative to 0:

$$\frac{\partial L_P}{\partial \mathbf{w}} = 0 \Rightarrow \mathbf{w} = \sum_{n=1}^N \alpha_n y_n \mathbf{x}_n$$

$$\frac{\partial L_P}{\partial b} = 0 \Rightarrow \sum_{n=1}^N \alpha_n y_n = 0$$

- Use the above results in the **Primal** Lagrangian L_P to get the **Dual** Lagrangian:

$$\underset{\mathbf{w}, b, \alpha}{\text{minimize}} \quad L_D(\mathbf{w}, b, \alpha) = \sum_{n=1}^N \alpha_n - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N \alpha_m \alpha_n y_m y_n (\mathbf{x}_m^\top \mathbf{x}_n)$$

$$\text{subject to} \quad \sum_{n=1}^N \alpha_n y_n = 0, \alpha_n \geq 0; n = 1, \dots, N.$$

Solving for α_n

- A **Quadratic Programming** problem in α
- Use “off-the-shelf” quadratic solvers for L_D
 - quadprog (MATLAB), CVXOPT
- Solution should satisfy certain conditions
- Also known as the **Karush-Kuhn-Tucker** (KKT) Conditions

The Karush-Kuhn-Tucker Conditions

$$\frac{\partial}{\partial \mathbf{w}} L_P(\mathbf{w}, b, \alpha) = \mathbf{w} - \sum_{n=1}^N \alpha_n y_n \mathbf{x}_n = 0 \quad (1)$$

$$\frac{\partial}{\partial b} L_P(\mathbf{w}, b, \alpha) = - \sum_{n=1}^N \alpha_n y_n = 0 \quad (2)$$

$$y_n \{\mathbf{w}^\top \mathbf{x}_n + b\} - 1 \geq 0 \quad (3)$$

$$\alpha_n \geq 0 \quad (4)$$

$$\alpha_n (y_n \{\mathbf{w}^\top \mathbf{x}_n + b\} - 1) = 0 \quad (5)$$

From Primal to Dual

- Using KKT conditions (1) and (2), we get:

$$\mathbf{w} = \sum_{n=1}^N \alpha_n y_n \mathbf{x}_n$$

$$0 = \sum_{n=1}^N \alpha_n y_n$$

- Now we only need to optimize using α_n 's
- This is done using the **dual formulation**

$$\begin{aligned} \underset{\mathbf{w}, b, \alpha}{\text{maximize}} \quad & L_D(\mathbf{w}, b, \alpha) = \sum_{n=1}^N \alpha_n - \frac{1}{2} \sum_{m,n=1}^N \alpha_m \alpha_n y_m y_n (\mathbf{x}_m^\top \mathbf{x}_n) \\ \text{subject to} \quad & \sum_{n=1}^N \alpha_n y_n = 0, \alpha_n \geq 0 \quad n = 1, \dots, N. \end{aligned}$$

Two Key Observations from Dual Formulation

Observation 1: Dot Product Formulation

- All training examples (\mathbf{x}_n 's) occur in *dot/inner products*
- Also recall the prediction using SVMs

$$\begin{aligned}y^* &= \text{sign}(\mathbf{w}^\top \mathbf{x}^* + b) \\&= \text{sign}\left(\left(\sum_{n=1}^N \alpha_n y_n \mathbf{x}_n\right)^\top \mathbf{x}^*\right) \\&= \text{sign}\left(\sum_{n=1}^N \alpha_n y_n \left(\mathbf{x}_n^\top \mathbf{x}^*\right)\right)\end{aligned}$$

- Replace the dot products with kernel functions
 - Kernel or non-linear SVM

Two Key Observations from Dual Formulation

Observation 2: Most α_n 's are 0

- KKT condition #5:


$$\alpha_n(y_n\{\mathbf{w}^\top \mathbf{x}_n + b\} - 1) = 0$$

- If \mathbf{x}_n **not** on margin

$$y_n\{\mathbf{w}^\top \mathbf{x}_n + b\} > 1$$

$$\Rightarrow \alpha_n = 0$$

- $\alpha_n \neq 0$ only for \mathbf{x}_n on margin
- These are the **support vectors**
- Only need these for prediction



What if data is not linearly separable?

- Cannot go for zero training error
- Still learn a maximum margin hyperplane
 - ① Allow some examples to be misclassified
 - ② Allow some examples to fall **inside** the margin
- How do you set up the optimization for SVM training


What if data is not linearly separable?

- Cannot go for zero training error
- Still learn a maximum margin hyperplane
 - ① Allow some examples to be misclassified
 - ② Allow some examples to fall **inside** the margin
- How do you set up the optimization for SVM training

What if data is not linearly separable?

- Cannot go for zero training error
- Still learn a maximum margin hyperplane
 - ① Allow some examples to be misclassified
 - ② Allow some examples to fall **inside** the margin
- How do you set up the optimization for SVM training

Cutting Some Slack



Introducing Slack Variables

- **Separable Case:** To ensure zero training loss, constraint was

$$y_n(\mathbf{w}^\top \mathbf{x}_n + b) \geq 1 \quad \forall n = 1 \dots N$$

- **Non-separable Case:** Relax the constraint

$$y_n(\mathbf{w}^\top \mathbf{x}_n + b) \geq 1 - \xi_n \quad \forall n = 1 \dots N$$

- ξ_n is called **slack variable** ($\xi_n \geq 0$)
- For misclassification, $\xi_n > 1$

Introducing Slack Variables

- **Separable Case:** To ensure zero training loss, constraint was

$$y_n(\mathbf{w}^\top \mathbf{x}_n + b) \geq 1 \quad \forall n = 1 \dots N$$

- **Non-separable Case:** Relax the constraint

$$y_n(\mathbf{w}^\top \mathbf{x}_n + b) \geq 1 - \xi_n \quad \forall n = 1 \dots N$$

- ξ_n is called **slack variable** ($\xi_n \geq 0$)
- For misclassification, $\xi_n > 1$

Relaxing the Constraint

- It is OK to have some misclassified training examples
 - Some ξ_n 's will be non-zero
- Minimize the number of such examples

- Minimize $\sum_{n=1}^N \xi_n$

- Optimization Problem for Non-Separable Case

$$\begin{aligned} &\underset{\mathbf{w}, b}{\text{maximize}} && f(\mathbf{w}, b) = \|\mathbf{w}\|^2 + C \sum_{n=1}^N \xi_n \\ &\text{subject to} && y_n(\mathbf{w}^\top \mathbf{x}_n + b) \geq 1 - \xi_n, \xi_n \geq 0 \quad n = 1, \dots, N. \end{aligned}$$

- C controls the impact of margin and the margin error.

Relaxing the Constraint

- It is OK to have some misclassified training examples
 - Some ξ_n 's will be non-zero
- Minimize the number of such examples

- Minimize $\sum_{n=1}^N \xi_n$

- Optimization Problem for Non-Separable Case

$$\begin{aligned} &\underset{\mathbf{w}, b}{\text{maximize}} && f(\mathbf{w}, b) = \|\mathbf{w}\|^2 + C \sum_{n=1}^N \xi_n \\ &\text{subject to} && y_n(\mathbf{w}^\top \mathbf{x}_n + b) \geq 1 - \xi_n, \xi_n \geq 0 \quad n = 1, \dots, N. \end{aligned}$$

- C controls the impact of margin and the margin error.

Estimating Weights

- What is the role of C ?
- Similar optimization procedure as for the separable case (QP for the dual)
- Weights have the same expression

$$\mathbf{w} = \sum_{n=1}^N \alpha_n y_n \mathbf{x}_n$$


- All training examples exist in dot products (*kernelizable*)
- Support vectors are slightly different
 - 1 Points on the margin ($\xi_n = 0$)
 - 2 Inside the margin but on the correct side ($0 < \xi_n < 1$)
 - 3 On the wrong side of the hyperplane ($\xi_n \geq 1$)

Outline


- 1 Supervised Learning
 - Learning Parameters
 - Risk Minimization for Parameter Learning
- 2 Grouping Learning Algorithms
 - 0-1 Loss
 - Approximation to 0-1 Loss
- 3 Maximum Margin Classifiers
 - Boundary with Maximum Margin
- 4 Support Vector Machines
 - SVM Learning
 - SVM for Non-Separable Case
 - Optimization Constraints
- 5 **More About Kernels**
 - Gaussian Kernel

Why Use Kernels?

- $x \in \mathbb{R}$
- No linear separator





- Map $x \rightarrow \{x, x^2\}$
- Separable in 2D space





Another Example

- $\mathbf{x} \in \mathbb{R}^2$
- No linear separator
- Map $\mathbf{x} \rightarrow \{x_1^2, \sqrt{2}x_1x_2, x_2^2\}$
- A circle as the decision boundary



Another Example

- $\mathbf{x} \in \mathbb{R}^2$
- No linear separator
- Map $\mathbf{x} \rightarrow \{x_1^2, \sqrt{2}x_1x_2, x_2^2\}$
- A circle as the decision boundary



The Gaussian Kernel

- The *squared dot product* kernel ($\mathbf{x}, \mathbf{x}' \in \mathbb{R}^2$):

$$k(\mathbf{x}, \mathbf{x}') = \mathbf{x}^\top \mathbf{x}' \triangleq \phi(\mathbf{x})^\top \phi(\mathbf{x}')$$

$$\phi(\mathbf{x}) = \{x_1^2, \sqrt{2}x_1x_2, x_2^2\}$$

- What about the Gaussian kernel (radial basis function)?

$$k(\mathbf{x}, \mathbf{x}') = \exp\left(-\frac{1}{2\sigma^2} \|\mathbf{x} - \mathbf{x}'\|^2\right)$$

Why is the Gaussian Kernel Mapping to Infinite Dimensions

- Assume $\sigma = 1$ and $\mathbf{x} \in \Re$ (denoted as x)

$$\begin{aligned} k(x, x') &= \exp(-x^2) \exp(-x'^2) \exp(2xx') \\ &= \exp(-x^2) \exp(-x'^2) \sum_{k=0}^{\infty} \frac{2^k x^k x'^k}{k!} \\ &= \sum_{k=0}^{\infty} \left(\frac{2^{k/2}}{\sqrt{k!}} x^k \exp(-x^2) \right) \left(\frac{2^{k/2}}{\sqrt{k!}} x'^k \exp(-x'^2) \right) \end{aligned}$$

- Using Maclaurin Series Expansion

$$k(x, x') = \begin{pmatrix} 1 \\ 2^{1/2} x^1 \exp(-x^2) \\ \frac{2^{2/2}}{2} x^2 \exp(-x^2) \\ \vdots \end{pmatrix} \times \begin{pmatrix} 1 \\ 2^{1/2} x'^1 \exp(-x'^2) \\ \frac{2^{2/2}}{2} x'^2 \exp(-x'^2) \\ \vdots \end{pmatrix}^{\top}$$

SVM Extensions

- ① Multiple classes
 - One vs. Rest
 - One vs. One
- ② One class SVM
- ③ Transductive SVM (Semi-supervised Learning)
- ④ Support Vector Regression
- ⑤ Custom kernels (Use a *Gram* matrix)

References